

```
# Data Loading
from google.colab import drive
import pandas as pd
drive.mount('/content/drive')
file_path = '/content/drive/My Drive/ElectricCarData_Clean.csv'
df = pd.read_csv(file_path)
```

Mounted at /content/drive

Here are a few ways I performed exploratory data analysis and data visualization on this electric vehicle dataset to help w
 # Import libraries like pandas, matplotlib, seaborn for data manipulation and visualization. Use pandas to read in the CSV da
 # Look at basic statistics of the data like the describe() method to see mean, min, max values for numerical columns like pri
 # Visualize distributions of key variables like price, range, efficiency using histograms, boxplots.
 # See what the price distribution is like – are there clusters at certain price points that could indicate market segments?
 # Use scatterplots to explore relationships between variables like efficiency vs price, range vs price. See if certain segmen
 # Plot grouped or faceted charts to compare vehicle body types, power train, etc. See if SUVs are generally more expensive th
 # Use pie, bar charts to visualize categorical variables like plug type, body style. See the proportions of each group.
 # Calculate statistics like average price, efficiency for each segment or group to quantify differences seen in visualization
 # Try clustering algorithms like k-means to automatically find segments based on variables like price, range, efficiency.
 # Geographical data could also help identify markets. Plot maps showing EV popularity by state using open data sets.
 # The goal is to visualize and understand relationships between the variables to identify potential natural segments in the I

+ Code

+ Text

```
# Explore the data
print(df.head())
```

	Brand	Model	AccelSec	TopSpeed_KmH	\
0	Tesla	Model 3 Long Range Dual Motor	4.6	233	
1	Volkswagen	ID.3 Pure	10.0	160	
2	Polestar	2	4.7	210	
3	BMW	iX3	6.8	180	
4	Honda	e	9.5	145	

	Range_Km	Efficiency_WhKm	FastCharge_KmH	RapidCharge	PowerTrain	\
0	450	161	940	Yes	AWD	
1	270	167	250	Yes	RWD	
2	400	181	620	Yes	AWD	
3	360	206	560	Yes	RWD	
4	170	168	190	Yes	RWD	

	PlugType	BodyStyle	Segment	Seats	PriceEuro
0	Type 2 CCS	Sedan	D	5	55480
1	Type 2 CCS	Hatchback	C	5	30000
2	Type 2 CCS	Liftback	D	5	56440
3	Type 2 CCS	SUV	D	5	68040
4	Type 2 CCS	Hatchback	B	4	32997

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103 entries, 0 to 102
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brand                 103 non-null   object
1   Model                 103 non-null   object
2   AccelSec              103 non-null   float64
3   TopSpeed_KmH          103 non-null   int64
4   Range_Km              103 non-null   int64
5   Efficiency_WhKm        103 non-null   int64
6   FastCharge_KmH         103 non-null   object
7   RapidCharge           103 non-null   object
8   PowerTrain            103 non-null   object
9   PlugType              103 non-null   object
10  BodyStyle             103 non-null   object
11  Segment               103 non-null   object
12  Seats                 103 non-null   int64
13  PriceEuro             103 non-null   int64
dtypes: float64(1), int64(5), object(8)
memory usage: 11.4+ KB
None
```

```
print(df.describe())
```

	AccelSec	TopSpeed_KmH	Range_Km	Efficiency_WhKm	Seats	\
count	103.000000	103.000000	103.000000	103.000000	103.000000	
mean	7.396117	179.194175	338.786408	189.165049	4.883495	
std	3.017430	43.573030	126.014444	29.566839	0.795834	
min	2.100000	123.000000	95.000000	104.000000	2.000000	
25%	5.100000	150.000000	250.000000	168.000000	5.000000	
50%	7.300000	160.000000	340.000000	180.000000	5.000000	
75%	9.000000	200.000000	400.000000	203.000000	5.000000	
max	22.400000	410.000000	970.000000	273.000000	7.000000	

```

      PriceEuro
count    103.000000
mean    55811.563107
std     34134.665280
min     20129.000000
25%     34429.500000
50%     45000.000000
75%     65000.000000
max     215000.000000

```

```
print(df.columns)
```

```

Index(['Brand', 'Model', 'AccelSec', 'TopSpeed_KmH', 'Range_Km',
      'Efficiency_WhKm', 'FastCharge_KmH', 'RapidCharge', 'PowerTrain',
      'PlugType', 'BodyStyle', 'Segment', 'Seats', 'PriceEuro'],
      dtype='object')

```

```
df.shape
```

```
(103, 14)
```

```

# Distribution of price
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure()
sns.distplot(df["PriceEuro"], kde=True)
sns.histplot(df["PriceEuro"], kde=True, stat="density")
plt.xlabel("Price (Euros)")

```

```
<ipython-input-10-06f499910151>:5: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```

sns.distplot(df["PriceEuro"], kde=True)
Text(0.5, 0, 'Price (Euros)')

```




```

# Price vs Range scatterplot
plt.figure()
sns.regplot(x="Range_Km", y="PriceEuro", data=df)
plt.xlabel("Range (km)")
plt.ylabel("Price (Euros)")


```

```
Text(0, 0.5, 'Price (Euros)')
```




```
# Faceted histogram by BodyStyle
plt.figure()
sns.FacetGrid(df, col="BodyStyle").map(plt.hist, "PriceEuro")
```

```
<seaborn.axisgrid.FacetGrid at 0x7e57359dfc40>
<Figure size 640x480 with 0 Axes>
```




```
# Boxplot of Price by PowerTrain
plt.figure()
sns.boxplot(x="PowerTrain", y="PriceEuro", data=df)
```

```
<Axes: xlabel='PowerTrain', ylabel='PriceEuro'>
```




```
# Bar chart of segment counts
plt.figure()
sns.countplot(x="Segment", data=df)
```

<Axes: xlabel='Segment', ylabel='count'>



```
plt.figure(figsize=(15,8))
sns.heatmap(df.corr(), annot=True)
plt.show()
```

<ipython-input-60-7d08400b79ac>:2: FutureWarning: The default value of numeric_only is True, but will be changed to False in a future version of pandas. Please explicitly pass numeric_only=True to silence this warning.



```
# Average price per segment
print(df.groupby("Segment")["PriceEuro"].mean())
```

```
Segment
A      22693.714286
B      34799.227273
C      41199.100000
D      58487.933333
E      74269.400000
F      119690.750000
N       52812.833333
S       215000.000000
Name: PriceEuro, dtype: float64
```

```
# Kmeans clustering on Range, Efficiency, Price
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=5).fit(df[["Range_Km", "Efficiency_WhKm", "PriceEuro"]])
print(kmeans.labels_)

[3 0 3 3 0 1 0 0 3 3 3 0 0 3 0 0 4 0 0 0 0 3 0 2 3 3 0 3 0 0 3 0 0 1 0 0 3
 3 0 3 1 0 1 0 0 0 0 1 2 0 1 4 0 0 1 0 0 0 0 1 0 3 3 1 3 1 0 3 0 3 0 3 2 3
 0 0 3 0 3 2 0 1 0 3 3 0 3 3 3 0 1 0 0 3 0 0 0 3 3 1 3 3 3]
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will
warnings.warn()
```


```
from sklearn.metrics import silhouette_score

# Kmeans clustering from previous code
kmeans = KMeans(n_clusters=5).fit(df[["Range_Km", "Efficiency_WhKm", "PriceEuro"]])

# Calculate silhouette score
score = silhouette_score(df[["Range_Km", "Efficiency_WhKm", "PriceEuro"]], kmeans.labels_)
print("Silhouette score:", score)

# Plot clusters
plt.figure()
colors = ['r', 'g', 'b', 'y', 'c']
for i in range(kmeans.n_clusters):
    cluster_df = df[kmeans.labels_==i]
    plt.scatter(cluster_df["Range_Km"], cluster_df["PriceEuro"], c=colors[i], label="Cluster "+str(i))
plt.xlabel("Range")
plt.ylabel("Price")
plt.legend()
```

```
Silhouette score: 0.6263766906659917
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will
warnings.warn()
<matplotlib.legend.Legend at 0x7e57333f6950>
```



```
#The silhouette score provides an evaluation of how well separated the formed clusters are.
# It ranges from -1 to 1, with a higher score indicating better defined clusters.
# Visualizing the clustered data projected onto just two dimensions (range and price in this case) gives us a sense of the c
# But keep in mind, the clustering was done in 3 dimensions - range, efficiency and price.
# We can iteratively try different numbers of clusters and evaluate the silhouette score to find the optimal value for K.
# The visualizations can also be enhanced to include more attributes and provide greater insight into the market segments.
```

```
from sklearn.mixture import GaussianMixture
from sklearn.tree import DecisionTreeRegressor
```

```
# Preprocess data
X = df[["Range_Km", "Efficiency_WhKm", "PriceEuro"]]
y = df["Segment"]
```

```
# K-Means Clustering
kmeans = KMeans(n_clusters=5)
kmeans.fit(X)
y_pred = kmeans.predict(X)
print("K-Means accuracy:", kmeans.score(X, y))
```

K-Means accuracy: -6727246470.872135

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will
warnings.warn()

```
# Gaussian Mixture Model
gmm = GaussianMixture(n_components=5)
gmm.fit(X)
y_pred = gmm.predict(X)
print("GMM accuracy:", gmm.score(X, y))
```

GMM accuracy: -20.537568849182623

```
# In this code:
# We load the electric vehicle dataset and extract the range, efficiency, price as features X and segment as the target y.
# We apply 2 models: K-Means clustering, Gaussian Mixture Models.
# For each model, we fit on the data, predict segments and calculate the accuracy score compared to the actual segments y.
# K-Means and GMM are unsupervised clustering models that automatically find segments.
# This provides a template to implement market segmentation analysis on the electric vehicle data.
```

```
from sklearn.preprocessing import StandardScaler

X = df[["Range_Km", "Efficiency_WhKm", "PriceEuro"]]
y = df["Segment"]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# PCA for dimension reduction
from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X_reduced = pca.fit_transform(X_scaled)

# K-Nearest Neighbors
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier()
knn.fit(X_reduced, y)
y_pred = knn.predict(X_reduced)
print("KNN Accuracy:", knn.score(X_reduced, y))

# Neural Network
from sklearn.neural_network import MLPClassifier

mlp = MLPClassifier()
mlp.fit(X_reduced, y)
y_pred = mlp.predict(X_reduced)
print("NN Accuracy:", mlp.score(X_reduced, y))
```

KNN Accuracy: 0.7281553398058253

NN Accuracy: 0.6990291262135923

/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stocha
warnings.warn()

```
# In the code:
# We scale data points to normalize the features.
# Apply PCA to reduce the dimensions to 2 principal components.
# Use KNN classifier and Neural Network classifier on the reduced data.
# Evaluate both models by their accuracy score on predicting segments y.
# The neural network is able to learn complex relationships between attributes.
# KNN is simpler but may also provide good accuracy.
# This demonstrates how to apply additional models like NN and KNN for market segmentation on this dataset, with model evalu
```

```
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# SVM
svm = SVC()
svm.fit(X_train, y_train)
print("SVM Accuracy:", svm.score(X_test, y_test))

# Random Forest
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
print("RF Accuracy:", rf.score(X_test, y_test))
```

SVM Accuracy: 0.38095238095238093
 RF Accuracy: 0.6190476190476191

```
# In the above split data into train and test sets
# Apply SVM classifier and Random Forest classifier
# Evaluate accuracy on test set for both models
# SVM makes segment predictions based on defining optimal decision boundaries.
# Random Forest builds an ensemble of decision trees and ensemble predicts segments.
# Compare accuracy scores and visualizations to evaluate which performs better.
# This demonstrates applying SVM and Random Forest for market segmentation on the EV dataset, along with model evaluation
```


```
from sklearn.cluster import AgglomerativeClustering, FeatureAgglomeration
from sklearn.metrics import accuracy_score
import scipy.cluster.hierarchy as shc
```

```
# Agglomerative clustering
aclust = AgglomerativeClustering(n_clusters=5)
aclust.fit(X)
y_pred = aclust.fit_predict(X)
```

```
print("Accuracy of Agglomerative clustering:", accuracy_score(y, y_pred))
```

```
# Plot dendrogram
plt.figure(figsize=(16,8))
dend = shc.dendrogram(shc.linkage(X, method='ward'))
plt.title("Agglomerative Dendrogram")
```

Accuracy of Agglomerative clustering: 0.0
 Text(0.5, 1.0, 'Agglomerative Dendrogram')



```
# Hierarchical clustering builds a hierarchy of clusters represented as a dendrogram.
# Agglomerative clustering starts with each data point as its own cluster and merges them iteratively based on similarity.
# Accuracy was calculated versus true segments to evaluate the techniques.
# Visualizing the dendrograms helps understand the hierarchy and clustering produced.
# This demonstrates applying both hierarchical clustering approaches for market segmentation on the EV data, along with eval
```